

ExoClick_Challenge_Solution

October 20, 2018

0.1 Exploratory Data Analysis

```
In [18]: import pandas as pd

        # reading csv as a dataframe
        df = pd.read_csv('ExoClick_challenge.csv')
        print('Shape of data:',df.shape)
        print('Shape of data if missing values are to be dropped:', df.dropna().shape)

('Shape of data:', (96223, 6))
('Shape of data if missing values are to be dropped:', (96213, 6))

In [19]: # Finding the unique values; Data Pre-Processing in case of redundancy;
        print('Unique countries:',df['country'].unique())
        print('Unique ad_type:',df['ad_type'].unique(), '\n')

        # Check for missing data
        print('Checking for attribute with missing values: \n',df.isna().any())

('Unique countries:', array(['EST', 'NCL', 'DZA', 'ISL', 'isl', 'dza', 'est', 'ncl'],
                             dtype=object))
('Unique ad_type:', array([10, 17,  2, 15]), '\n')
('Checking for attribute with missing values: \n', datetime        False
client                False
country               False
ad_type               False
impressions           False
value                 True
dtype: bool)

In [20]: # dropping missing values as insignificant number of them are present
        df = df.dropna()
        print('Dataframe view after removing missing values\n',df.head(5))

        # rows with redundant country codes
        print('\n','Rows with redundant Country Codes')
```

```

print(df[df['country']=='isl'])
print(df[df['country']=='dza'])
print(df[df['country']=='est'])
print(df[df['country']=='ncl'])

('Dataframe view after removing missing values\n',
0 2016-06-06*06:00:00      1    EST      10      137    0.1370
1 2016-06-12*22:00:00      2    NCL      10     4424   91.7207
2 2016-06-05*20:00:00      1    DZA      10      271    0.2710
3 2016-06-23*17:00:00      1    EST      17         1    0.0010
4 2016-06-02*10:00:00      3    DZA      10     11701   4.8719)

('\n', 'Rows with redundant Country Codes')
      datetime client country ad_type impressions  value
624  2016-06-16*07:00:00      9    isl         2         89  0.4894
6397 2016-06-03*23:00:00     25    isl         2        129  0.6450
15119 2016-06-11*10:00:00      3    isl        17         7  0.0049
30699 2016-06-01*14:00:00     13    isl         2         94  0.0000
      datetime client country ad_type impressions  value
768  2016-06-25*14:00:00     23    dza         2        359  4.7194
79898 2016-06-29*06:00:00      3    dza        10       1439  0.8040
      datetime client country ad_type impressions  value
17855 2016-06-03*05:00:00     38    est         2         34  0.3554
27891 2016-06-06*10:00:00     52    est         2          5  0.0250
61551 2016-06-28*17:00:00      9    est         2        580  2.9085
      datetime client country ad_type impressions  value
61833 2016-06-09*09:00:00      5    ncl        10        101  1.01

```

```

In [21]: # converting the countries with lower case to upper case
row_numbers_to_process = []
row_numbers_to_process.extend(df[df['country']=='isl'].index)
row_numbers_to_process.extend(df[df['country']=='dza'].index)
row_numbers_to_process.extend(df[df['country']=='est'].index)
row_numbers_to_process.extend(df[df['country']=='ncl'].index)

for num in row_numbers_to_process:
    df.loc[num,'country'] = df.loc[num,'country'].upper()

print('Unique Countries after Processing:', df['country'].unique())

('Unique Countries after Processing:', array(['EST', 'NCL', 'DZA', 'ISL'], dtype=object))

In [22]: # computing summary statistics
print('Shape of Dataset',df.shape)
print(df.describe())

```

```


('Shape of Dataset', (96213, 6))
      client      ad_type  impressions      value

```

count	96213.000000	96213.000000	96213.000000	96213.000000
mean	16.665004	6.642190	4045.791535	4.002480
std	14.327285	5.593171	31902.587268	15.084129
min	1.000000	2.000000	1.000000	0.000000
25%	4.000000	2.000000	7.000000	0.020100
50%	13.000000	2.000000	42.000000	0.140000
75%	25.000000	10.000000	379.000000	1.104600
max	73.000000	17.000000	941289.000000	846.000000


```
In [78]: print('Frequency of ad_types\n')
         df['ad_type'].hist();
```

Frequency of ad_types



```
In [81]: print('Country frequency')
         df['country'].hist();
```

Country frequency



0.2 Computing eCPM

```
In [82]: # computing eCPM for each segment given by a country and ad_type
ecpm_df = df.groupby(['country', 'ad_type']).sum()
ecpm_df['eCPM'] = ''
ecpm_df['eCPM'] = (ecpm_df['value']*1000)/(ecpm_df['impressions'])
print(ecpm_df[['impressions', 'value', 'eCPM']])
```

		impressions	value	eCPM
country	ad_type			
DZA	2	209732533	64402.8197	0.307071
	10	99589395	49136.4866	0.493391
	15	36051	48.1423	1.335394
	17	288743	130.5920	0.452278
EST	2	30512420	85804.5274	2.812118
	10	14598906	40917.6501	2.802789
	15	22860	23.7174	1.037507
	17	83608	199.4982	2.386114
ISL	2	15694168	14342.6641	0.913885
	10	6366502	12074.8048	1.896615
	15	5117	6.6349	1.296639
	17	58403	104.9812	1.797531
NCL	2	7764619	66507.4754	8.565453
	10	4481430	51331.6085	11.454292
	15	5546	16.7132	3.013559

17 17440 42.2456 2.422339

This clearly allows us to infer that 'NCL' is the country with highest eCPM for the ad_type 10. We can also draw some other conclusions such as- - In almost all countries, ad_type 2 is having the highest impressions followed by ad_type 10 - The country 'NCL' shows significantly higher levels of eCPM values for all values of ad_type

```
In [24]: # creating a dataframe with separate rows for date and time for ease of dataprocessing
date_time = pd.DataFrame(df['datetime'].str.split('*',1).tolist(),
                          columns = ['date', 'time'])

df['date'] = date_time['date']
df['time'] = date_time['time']
df = df.drop(['datetime'], axis=1)
df = df[['date', 'time', 'client', 'country', 'ad_type', 'impressions', 'value']]
df.head(5)
```

```
Out[24]:
```

	date	time	client	country	ad_type	impressions	value
0	2016-06-06	06:00:00	1	EST	10	137	0.1370
1	2016-06-12	22:00:00	2	NCL	10	4424	91.7207
2	2016-06-05	20:00:00	1	DZA	10	271	0.2710
3	2016-06-23	17:00:00	1	EST	17	1	0.0010
4	2016-06-02	10:00:00	3	DZA	10	11701	4.8719

0.3 Daily Time Series


```
In [25]: import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [60]: # hourly time series including the dates
NCL_df = df[df['country']=='NCL']
NCL_df['date'] = pd.to_datetime(NCL_df['date'], format='%Y-%m-%d')
NCL_df = NCL_df.drop(['time'], axis=1)
NCL_df = NCL_df[['date', 'client', 'country', 'ad_type', 'impressions',
                  'value']]


NCL_df = NCL_df.sort_values(by= ['date'])
NCL_df = NCL_df.dropna()
NCL_df = NCL_df.groupby(['ad_type', 'date']).sum()
NCL_df['eCPM'] = ''
NCL_df['eCPM'] = (NCL_df['value']*1000)/(NCL_df['impressions'])
NCL_df = NCL_df.reset_index()
NCL_df.set_index('date', inplace=True)
```

```
In [61]: # dataframes based on specific ad_types
NCL_02_df = NCL_df[NCL_df['ad_type']== 2]
NCL_10_df = NCL_df[NCL_df['ad_type']==10]
NCL_15_df = NCL_df[NCL_df['ad_type']==15]
NCL_17_df = NCL_df[NCL_df['ad_type']==17]
```


```
In [55]: NCL_02_df[['eCPM']].plot(figsize=(15,3), linewidth=2, fontsize=15)
plt.xlabel('date', fontsize=15);
plt.title('Advertisement Type - 2: Variation of eCPM across July, 2016', fontsize=15)
```




```
In [56]: NCL_10_df[['eCPM']].plot(figsize=(15,3), linewidth=2, fontsize=15)
plt.xlabel('date', fontsize=15);
plt.title('Advertisement Type - 10: Variation of eCPM across June, 2016', fontsize=15)
```



```
In [57]: NCL_15_df[['eCPM']].plot(figsize=(15,3), linewidth=2, fontsize=15)
plt.xlabel('date', fontsize=15);
plt.title('Advertisement Type - 15: Variation of eCPM across June, 2016', fontsize=15)
```



```
In [62]: NCL_17_df[['eCPM']].plot(figsize=(15,3), linewidth=2, fontsize=15)
plt.xlabel('date', fontsize=15);
plt.title('Advertisement Type - 17: Variation of eCPM across June, 2016', fontsize=15)
```



Clearly, we can draw certain conclusions from the information received so far from the specific ad_type time series plots regarding the variation of the metric eCMP. Some of them are listed below- * Advertisement Type 2- The highest value of eCPM occurs on 13th of June, 2016. Some relatively high values of the eCPM are also observed on 10th of June, 2016 and 18th of June, 2016. * Advertisement Type 10- The highest value of eCPM occurs on 2nd of June, 2016 * Advertisement Type 15- It has a more seasonal behaviour with crests and troughs occurring routinely. The highest value of eCPM is attained on 14th of June and 26th of June, 2016 * Advertisement Type 17- We can infer that the highest cost of eCPM arises on 18th of June, 2016 which continues to go through smaller crests and falls. The other important high values of eCPM occur on 3rd of June, 7th of June, 22nd of June, 25th of June

0.4 Hourly Time Series

```
In [33]: import warnings
import matplotlib.pyplot as plt
warnings.filterwarnings('ignore')

In [34]: # hourly time series excluding the dates
DZA_df = df[df['country']=='DZA']
DZA_df['time'] = pd.to_datetime(DZA_df['time'], format='%H:%M:%S').dt.hour
DZA_df = DZA_df.drop(['date'], axis=1)
DZA_df = DZA_df[['time', 'client', 'country', 'ad_type', 'impressions',
                  'value']]
DZA_df = DZA_df.sort_values(by= ['time'])
DZA_df = DZA_df.dropna()
DZA_df = DZA_df.groupby(['ad_type', 'time']).sum()
DZA_df['eCPM'] = ''
DZA_df['eCPM'] = (DZA_df['value']*1000)/(DZA_df['impressions'])
```

```
DZA_df = DZA_df.reset_index()
DZA_df.set_index('time', inplace=True)
```


```
In [35]: # dataframes based on ad_types
```

```
DZA_02_df = DZA_df[DZA_df['ad_type']== 2]
DZA_10_df = DZA_df[DZA_df['ad_type']==10]
DZA_15_df = DZA_df[DZA_df['ad_type']==15]
DZA_17_df = DZA_df[DZA_df['ad_type']==17]
```

```
In [36]: DZA_02_df[['eCPM']].plot(figsize=(15,3), linewidth=2, fontsize=15)
```

```
plt.xlabel('time', fontsize=15);
```


```
plt.title('Advertisement Type - 2: Variation of eCPM across hours (00-23)', fontsize=15);
```



```
In [37]: DZA_10_df[['eCPM']].plot(figsize=(15,3), linewidth=2, fontsize=15)
```

```
plt.xlabel('time', fontsize=15);
```


```
plt.title('Advertisement Type - 10: Variation of eCPM across hours (00-23)', fontsize=15);
```




```
In [38]: DZA_15_df[['eCPM']].plot(figsize=(15,3), linewidth=2, fontsize=15)
```

```
plt.xlabel('time', fontsize=15);
```

```
plt.title('Advertisement Type - 15: Variation of eCPM across hours (00-23)', fontsize=15);
```

```
In [39]: DZA_17_df[['eCPM']].plot(figsize=(15,3), linewidth=2, fontsize=15)
plt.xlabel('time', fontsize=15);
plt.title('Advertisement Type - 17: Variation of eCPM across hours (00-23)', fontsize=15)
```



Clearly, we can draw certain conclusions from the information received so far from the specific ad_type time series plots regarding the variation of the metric eCPM. Some of them are listed below- * Advertisement Type 2- The highest value of eCPM is attained at 0200 hours in night. Some relatively high values of the eCPM are also achieved at 1400 hours and 2100 hours. * Advertisement Type 10- The highest value of eCPM is achieved at 1100 hours in the morning with comparatively smaller yet significantly high values attained also at 0600 hours in morning and 2100 hours in night * Advertisement Type 15- The highest value of eCPM is attained at 1300 hours in afternoon and 2100 hours in night * Advertisement Type 17- We can infer that the highest cost of eCPM arises at 0100 hours in night which continues to go through smaller crests and falls till it again reaches a significant high value at around 1400 hours in the afternoon.

Now, we have decided to generate an overview of how the time series will look if both the date and time as unique factors will be retained and the dataframe sorted based on both the factors. The four plots following the 2 sections of code have implemented this feature

```
In [49]: # hourly time series including the dates
DZA_df = df[df['country']=='DZA']
DZA_df['datetime'] = pd.to_datetime(DZA_df['date'] + ' ' + DZA_df['time'])
DZA_df = DZA_df.drop(['date','time'], axis=1)
DZA_df = DZA_df[['datetime', 'client', 'country', 'ad_type', 'impressions',
                  'value']]
DZA_df = DZA_df.sort_values(by= ['datetime'])
```

```

DZA_df = DZA_df.dropna()
DZA_df = DZA_df.groupby(['ad_type', 'datetime']).sum()
DZA_df['eCPM'] = ''
DZA_df['eCPM'] = (DZA_df['value']*1000)/(DZA_df['impressions'])
DZA_df = DZA_df.reset_index()
DZA_df.set_index('datetime', inplace=True)

```

In [50]: # dataframes based on ad_types

```


DZA_02_df = DZA_df[DZA_df['ad_type']== 2]
DZA_10_df = DZA_df[DZA_df['ad_type']==10]
DZA_15_df = DZA_df[DZA_df['ad_type']==15]
DZA_17_df = DZA_df[DZA_df['ad_type']==17]

```

In [51]: DZA_02_df[['eCPM']].plot(figsize=(15,3), linewidth=1, fontsize=15)

```
plt.xlabel('datetime', fontsize=15);
```


```
plt.title('Advertisement Type - 2: Hourly Variation of eCPM across June, 2016', fontsize=15);
```




In [43]: DZA_10_df[['eCPM']].plot(figsize=(15,3), linewidth=1, fontsize=15)

```
plt.xlabel('datetime', fontsize=15);
```


```
plt.title('Advertisement Type - 10: Hourly Variation of eCPM across June, 2016', fontsize=15);
```



```
In [44]: DZA_15_df[['eCPM']].plot(figsize=(15,3), linewidth=1, fontsize=15)
plt.xlabel('datetime', fontsize=15);
plt.title('Advertisement Type - 15: Variation of eCPM across June, 2016', fontsize=15);
```



```
In [45]: DZA_17_df[['eCPM']].plot(figsize=(15,3), linewidth=1, fontsize=15)
plt.xlabel('datetime', fontsize=15);
plt.title('Advertisement Type - 17: Variation of eCPM across June, 2016', fontsize=15);
```




0.5 Forecasting using a Time Series Model

In this segment, we are demonstrating the idea that we can create forecasts for eCPM using the `ad_type` specific dataframes created for a particular country. In total, we may need 16 time series models to accurately capture the eCPM values for each of the 4 specific `ad_type` and 4 specific countries. We are demonstrating for NCL with `ad_type` 2.

```
In [111]: import numpy as np
from statsmodels.tsa.arima_model import ARIMA
from statsmodels.graphics.tsaplots import plot_acf
from statsmodels.graphics.tsaplots import plot_pacf
from math import sqrt
from matplotlib import pyplot
```


```
In [148]: # visualisation
series = NCL_02_df[['eCPM']]
series.plot(figsize=(15,3))
# univariate series
print('Variance of original series:',series.var());
print('Variance of differenced series',series.diff().var());
```

```
('Variance of original series:', eCPM    1.421153
dtype: float64)
('Variance of differenced series', eCPM    2.390404
dtype: float64)
```



As we can clearly see, the series do not have any visible trend. A pattern of seasonality may appear but visibly, it doesn't seem so. Further, the variance of the series also looks pretty fine with little increase as the series progresses. However, at this stage, we will restrict ourselves from taking a logarithmic or square root transformation.

```
In [146]: # examination of acf and pacf
pyplot.figure()
pyplot.subplot(211)
plot_acf(series, ax=pyplot.gca())
pyplot.subplot(212)
plot_pacf(series, ax=pyplot.gca())
pyplot.show()
```



Both ACF and PACF are not showing any significant lags. We may start with $p=0$, $q=0$. Although, we may directly plug in these values in an ARIMA model but it would be better if we perform grid search to further develop some idea.

```
In [121]: # Grid search to confirm the chosen model parameters
# evaluate an ARIMA model for a given order (p,d,q) and return RMSE
def evaluate_arima_model(X, arima_order):
    # prepare training dataset
    X = X.astype('float32')
    train_size = int(len(X) * 0.50)
    train, test = X[0:train_size], X[train_size:]
    history = [x for x in train]
    # make predictions
    predictions = list()
    for t in range(len(test)):
        model = ARIMA(history, order=arima_order)
        # model_fit = model.fit(dispatch=0)
        model_fit = model.fit(trend='nc', disp=0)
        yhat = model_fit.forecast()[0]
        predictions.append(yhat)
        history.append(test[t])
    # calculate out of sample error
    mse = mean_squared_error(test, predictions)
    rmse = sqrt(mse)
    return rmse
```

```

# evaluate combinations of p, d and q values for an ARIMA model
def evaluate_models(dataset, p_values, d_values, q_values):
    dataset = dataset.astype('float32')
    best_score, best_cfg = float("inf"), None
    for p in p_values:
        for d in d_values:
            for q in q_values:
                order = (p,d,q)
                try:
                    mse = evaluate_arima_model(dataset, order)
                    if mse < best_score:
                        best_score, best_cfg = mse, order
                        print('ARIMA%s RMSE=%.3f' % (order,mse))
                except:
                    continue
    print('Best ARIMA%s RMSE=%.3f' % (best_cfg, best_score))

# evaluate parameters
p_values = range(0, 5)
d_values = range(0, 3)
q_values = range(0, 5)
warnings.filterwarnings("ignore")
evaluate_models(series.values, p_values, d_values, q_values)

```

```

ARIMA(0, 0, 1) RMSE=4.836
ARIMA(0, 0, 2) RMSE=4.231
ARIMA(0, 0, 3) RMSE=2.741
ARIMA(0, 1, 1) RMSE=1.469
ARIMA(0, 2, 1) RMSE=1.740
ARIMA(1, 0, 0) RMSE=1.710
ARIMA(1, 1, 0) RMSE=1.737
ARIMA(1, 2, 0) RMSE=2.433
ARIMA(2, 1, 0) RMSE=1.736
ARIMA(2, 2, 0) RMSE=2.421
ARIMA(3, 1, 0) RMSE=1.696
ARIMA(4, 0, 0) RMSE=2.276
ARIMA(4, 1, 0) RMSE=1.634
Best ARIMA(0, 1, 1) RMSE=1.469

```

```
In [138]: warnings.filterwarnings('ignore')
```

```

# intentionally training on first 25 days and testing on last 15 days due to lack of
train, test = series[['eCPM']].values[0:25], series[['eCPM']].values[15:29]

# fit model
model = ARIMA(train, order=(0,1,1))

```

```
model_fit = model.fit(dis=0)
print(model_fit.summary())
```

ARIMA Model Results

```
=====
Dep. Variable:          D.y    No. Observations:          24
Model:                ARIMA(0, 1, 1)    Log Likelihood          -39.667
Method:                css-mle    S.D. of innovations          1.239
Date:                  Sat, 20 Oct 2018    AIC          85.334
Time:                  14:03:49    BIC          88.868
Sample:                1    HQIC          86.271
=====
```


```
=====
              coef    std err          z      P>|z|      [0.025      0.975]
-----
const         -0.0321     0.066     -0.487     0.631     -0.161     0.097
ma.L1.D.y     -0.7814     0.227    -3.446     0.002     -1.226    -0.337
=====
```


Roots

```
=====
              Real          Imaginary          Modulus          Frequency
-----
MA.1         1.2797          +0.0000j          1.2797          0.0000
=====
```

In [142]: # plot residual errors

```
residuals = DataFrame(model_fit.resid)
residuals.plot(figsize=(15,3))
pyplot.show()
residuals.plot(kind='kde', figsize=(15,3))
pyplot.show()
print(residuals.describe())
```





```

0
count    24.000000
mean      0.114608
std       1.270540
min       -2.817981
25%      -0.355763
50%       0.244304
75%       1.093785
max        2.076164

```

The residuals, although not exactly following normality as we may expect, however, they are somewhat resembling.

```

In [122]: warnings.filterwarnings('ignore')
          from sklearn.metrics import mean_squared_error

          history = [x for x in train]
          predictions = []
          for t in range(len(test)):
              model = ARIMA(history, order=(0,1,1))
              model_fit = model.fit(dispatch=0)
              output = model_fit.forecast()
              yhat = output[0]
              predictions.extend(yhat)
              obs = test[t]
              history.append(obs)
              print('predicted=%f, expected=%f' % (yhat, obs))
          error = mean_squared_error(test, predictions)
          print('Test MSE: %.3f' % error)

```

```

predicted=7.798187, expected=6.773410
predicted=7.280596, expected=8.850637
predicted=7.853268, expected=10.087144
predicted=8.271701, expected=8.746229
predicted=8.355729, expected=7.046965
predicted=8.211132, expected=8.770722
predicted=8.201925, expected=9.858574

```




```
predicted=8.352102, expected=7.629319
predicted=8.310149, expected=8.561703
predicted=8.293146, expected=5.628232
predicted=8.045681, expected=7.637763
predicted=7.915104, expected=9.602250
predicted=8.014216, expected=7.838834
predicted=8.012606, expected=7.195957
Test MSE: 1.779
```

```
In [123]: # compute predictions dataframe
predictions = pd.DataFrame({'predictions':predictions})
true = [true for true in series['eCPM'][15:29].values]
predictions['dates'] = date
predictions['true'] = true
predictions = predictions[['dates', 'true', 'predictions']]
predictions.set_index('dates', inplace=True)
```

```
In [124]: # plot predictions
predictions.plot(figsize=(15,5))
```

```
Out[124]: <matplotlib.axes._subplots.AxesSubplot at 0x7fa4317515d0>
```



Under this analysis, our idea was to develop one model for each of the advertisement_types for each of the countries as they seems to be following significantly different behaviour from each other as exhibited from their previous plots.

Although, our prediction model seems to be not significantly precise in terms of forecasting the eCPM metric but it has successfully captured the behaviour in an approximate manner. The values seems to be showing an average of the crests and troughs. A more complex model may be able to capture the specific features but a better remedy would be to use several months data and compare the trend and seasonality accordingly. Such metrics often see fluctuations as observed across a year.

LSTM models may also be employed to fit the time series data. However, again, such models need extensively large amount of data.